
The BACON system for equation discovery from scientific data: Transforming classical artificial intelligence with modern machine learning approaches

Jonah Miller
University of Cambridge
jm2399@cantab.ac.uk

Soumya Banerjee
University of Cambridge
sb2333@cam.ac.uk

Abstract

BACON is a heuristic-based computational scientific discovery system, which aims to find invariants in multivariable systems. We rebuilt BACON in a modern computing language, and we improve the noise-resilience of BACON. We demonstrate how such classical AI systems can be understandable, yet powerful. We applied our framework to a number of exemplar problems in physics and mathematics. Our BACON also outperformed PySR - a modern method utilising symbolic regression on a neural network - conclusively in specific environments on small datasets.

We suggest that there is potential in these forgotten approaches that modern deep learning systems can learn from. Integrative approaches that combine heuristic approaches like BACON with modern deep learning can be very helpful. We suggest integrating modern deep learning approaches and large-language models with heuristic-based classical AI approaches as a way to analyse large scientific datasets.

Introduction

Discovering the numeric laws that shape experimental observations is a time-consuming endeavour filled with trial and error. Attempts at this are seen as early as the 16th century when astronomer Johannes Kepler deduced that planetary motion was an ellipse only after years of studiously trying to understand the observations of fellow astronomer Tycho Brahe. In the current age of computers this task can be deferred to programs which can spot patterns, relations and invariants in data, whilst performing the task faster and more accurately than any human. The initial program pioneering this field was BACON [8].

Fast-forward to the modern day, deep learning algorithms form the recent research in this area. Classical AI techniques like those used in BACON have been forgotten. Modern deep neural networks (DNNs) lack the explainability that BACON offers. Another issue is the time and computational complexity taken to train DNNs is immense in comparison to heuristic-based classical AI methods such as BACON.

Methods

The core functionality of BACON . 1 is based on one heuristic applied repeatedly. To find the relationship between two variables X and Y , if X increases whilst Y decreases (or vice versa), consider their product XY . If both increase or decrease, then consider the divisor $\frac{X}{Y}$. Table 1 demonstrates the path BACON . 1 takes to uncover Kepler’s 3rd law for synthetic variables D and P . Langley’s explanation of why this works is that BACON . 1 is a trend detector [11]. It systematically analyses and ranks the trends found, which allows it to uncover further patterns leading to the true relationship. The algorithms used to compare two variables are referred to as the Space of Laws. Langley also

introduced another solution to the Space of Laws – BACON . 6. It uses correlation coefficients to rank the most accurate form of an equation, given that the general relation (without coefficients) is already known.

Planet	Distance (D)	Period (P)	$\frac{D}{P}$	$\frac{D^2}{P}$	$\frac{D^2}{P^2}$	$\frac{D^3}{P^2}$
A	1.0	1.0	1.0	1.0	1.0	1.0
B	4.0	8.0	0.5	2.0	0.25	1.0
C	9.0	27.0	0.333	3.0	0.111	1.0

Table 1: An example of the BACON . 1 algorithm discovering Kepler’s 3rd law from a noiseless planetary system. The program is acting on 3 different synthetic planets which obey the same laws of motion (data from [10]).

BACON can also function with multiple variables and this functionality was first show in BACON . 3. It automatically sorts the data into a tree with the branches between representing different variables, with the dependent variable as the lowest layer. The program iterates through the layers in the tree plugging them into BACON . 1 to form a new relationship for the dependent variable. A demonstration of how this works is in Figure 3. BACON . 5 was an advance on BACON . 3 considering symmetry in the tree it cherrypicks branches to use, reducing the complexity of iteration. The space of functions used to iterate through the tree is referred to as the Space of Data.

The DNN used as comparison in this project is PySR [3]. PySR uses symbolic regression to find invariants. Symbolic regression works by imagining the search space as possible mathematical expressions. The expressions are narrowed down to the most accurate and often have a bound on complexity and simplicity in the final result. PySR is also a much more versatile and powerful program with customisability and the ability to run in large and complex search spaces where the only bound is the computational power. As such, it can handle invariants such as in differential equations and discontinuous environments. BACON cannot compete. However there are downsides. As it is trained on a deep neural network [4] there is no interface or explainability giving a reason for its outputs. This reduces the trustworthiness of the model.

The final product, BACON . 7, uses BACON . 1 as the Space of Laws and various layer-based approaches to traverse the Space of Data. These work by starting at the lower level of the tree and recursing upward. The expressions found at each layer are ranked with one chosen as truth, and then symmetrically applied. The output is then averaged, and a tree is formed. This is done until the tree has been fully manoeuvred. A visualisation of this in Figure 3. Multiple novel approaches were used to select the expression at each layer. Additional details regarding this are available in the Appendix. The Appendix also details other experiments such as a Monte Carlo Tree Search (MCTS) method overlaid on the BACON framework to choose the correct expression.

We release open-source code which will aid further development: <https://github.com/JonahMiller/BACON>.

Results

Criterion for success

Langley’s approach to testing prescribes that BACON can find an invariant if one combination of the hyperparameters inputted into the system allows the detection of the invariant. A similar metric is used for determining if PySR can find the equation; the correct form has to appear in its list of found equations. For α and β unknown constants the corrects forms to be found are:

$$V_{\text{BACON}} = \frac{M(\alpha T + \beta)}{P}, \quad I_{\text{BACON}} = \frac{\alpha T D^2}{(L + \beta)} \quad (1)$$

V_{BACON} are BACON’s predictions for V when fed the noisy dataset. V_{PySR} is correspondingly used later for PySR’s predictions.

The Ideal Gas Law

Figure 1 is a graph of the MSE differences between the equations that BACON . 7 and PySR predict from inputted noisy data ($V + n$), against the true data (V) and the noisy data. There’s a seemingly linear increase in the log of the MSE found as noise increases in all aspects. On the noisy data BACON . 7 is less accurate than PySR. It implies that PySR overfits on the dataset. This is verified when




Figure 1: For $V = \frac{M(T+273)}{P}$ and denoting the volume with noise added as $V + n$, the graph demonstrates the MSE between predicted models from BACON.7 and PySR with $V + n$ and V . The MSE with $V + n$ is how well the model predicts the data it is trained on (solid lines). The MSE with V describes how it predicts the true data with BACON.7 creating a model *an order of magnitude more accurate than PySR*. We hypothesise this is due to BACON.7's averaging through the Space of Data incidentally cancelling Gaussian noise, whereas PySR is trained to overfit on the data as it has no prior knowledge of there being noise. The latter explains as well why PySR displays a better model when compared with $V + n$ than BACON.

comparing V_{PySR} with $V + n$ where its model is more accurate than BACON.7. BACON.7 does a better job of deducing the true V in a noisy environment. Moreover, it does this consistently over all the noise percentages tested on. *On this dataset BACON.7 outperforms PySR by creating a model that is an order of magnitude more accurate through its superior ability to remove the noise.*¹

Ohm's Law

The results on Ohm's laws are displayed in Figure 2. They display similar trends for BACON.7. PySR could not find the correct equation form apart from for 0% noise so instead the results from the best prediction PySR returns is shown. PySR is consistently at least an order of magnitude worse

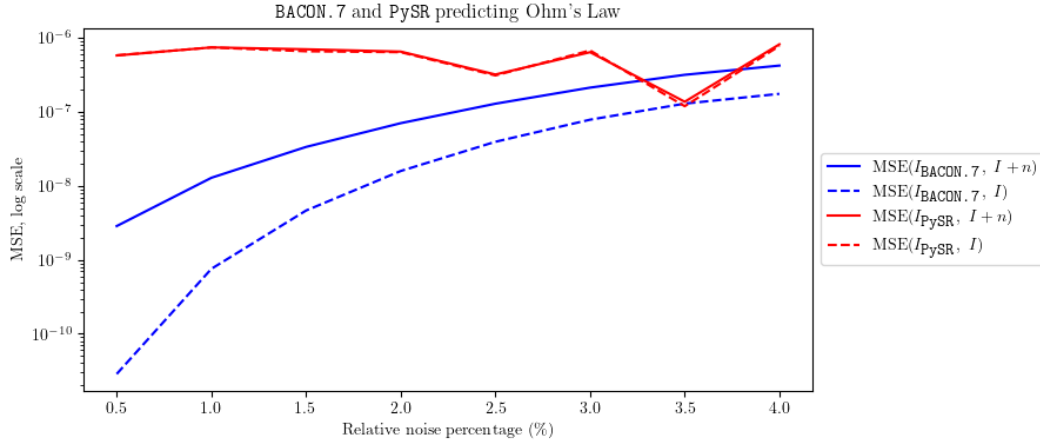


Figure 2: The results for Ohm's Law when using the same methodology as the Ideal Gas Law in Figure 1. BACON.7 again is better at predicting the I than $I + n$. Here PySR is not able to deduce the correct form of Ohm's Law so the results are taken from the best prediction it gives. In almost each case *PySR's predictions are worse than BACON.7's against both the noisy and true data*. The likely reason is that the correct form of the equation cannot be inferred by PySR from only 27 datapoints as PySR was built to function efficiently on larger, more complicated datasets. BACON.7 works best at this scale, where it strips out noise through averaging.

¹A more detailed discussion analysing the results can be found in the Appendix F.7.

in accuracy, than BACON.7 on I . Also, the difference between the noisy and noiseless forms for PySR is imperceptible. BACON.7 finds the correct form and as before, improves when run against the noiseless data. BACON.7 *again produces a model that strips away the noise more accurately than* PySR. *Moreover* PySR *cannot even deduce the correct form of the equation*.

Ramanujan birthday problem

We also applied our BACON.6 reconstruction to the Birthday Problem. This is demonstrated through the Birthday Problem: in a year of $n \geq 1$ days, what is the minimal number of people in a room such that the probability of at least two sharing a birthday is over 50%. This is extrapolated to, assuming each person enters the room sequentially, what is the expected number of people to enter, for a birthday to first be shared. Ramanujan discovered the answer is $\lfloor Q(n) + 1 \rfloor$ [1], where

$$Q(n) = \sum_{k=1}^n \frac{n!}{(n-k)! n^k} \quad (2)$$

$$= \sqrt{\frac{\pi n}{2}} - \frac{1}{3} + \frac{1}{12} \sqrt{\frac{\pi}{2n}} - \frac{4}{135n} + \mathcal{O}\left(\frac{1}{n^{\frac{3}{2}}}\right) \quad (3)$$

$$\approx 1.253\sqrt{n} - 0.333 + \frac{0.104}{\sqrt{n}} - \frac{0.030}{n} \quad (4)$$

When given an appropriate answer form without coefficients to BACON.6² with datapoints for $n \in \{2, 3, 4, 5, 6, 7, 8, 9\}$ and associated $Q(n)$, it discovers that:

$$Q_{\text{BACON.6}} = 1.252\sqrt{n} - 0.326 + \frac{0.086}{\sqrt{n}} - \frac{0.012}{n} \quad (6)$$

The MSE between $Q_{\text{BACON.6}}$ and the real $Q(n)$ is 5.4×10^{-9} . PySR comparatively struggles. Given the same data, and the *binary_operators* from Figure 15 expanded to include powers, the best result it finds is:

$$Q_{\text{PySR}} = n^{0.557} + \frac{0.324 \frac{5.28}{n}}{n} \quad (7)$$

This has MSE 3.8×10^{-5} to the true form – 4 orders of magnitude worse than BACON.6. This further lends credibility to the idea seen throughout that PySR does not function well on few datapoints. This also demonstrates the strength of BACON’s heuristic mechanisms to be incredibly accurate in an environment where the state-of-the-art performs poorly.

Summary and conclusions

For smaller datasets, BACON consistently thrives whilst PySR struggles. BACON is made for this environment, whereas PySR is made to overfit on complicated, large datasets whilst applying their biases towards simplicity. When approaching this threshold, it is PySR that thrives whilst BACON suffers (see experiments in the Appendix on Black’s Law). Here is an - albeit niche - situation where classical methods outperform modern techniques. It amplifies the need to reproduce, study and understand these seemingly anachronistic mechanisms and see what lessons can be taken going forward.

We applied this framework to a number of exemplar problems in physics and mathematics. Our results suggest that BACON is good at reducing noise and inferring the correct equation in smaller datasets, whereas PySR is significantly more successful on larger, noisier, datasets.

The broad goal of this research project was to combine modern approaches to AI with the classical. Both have strengths which, when efficiently combined, could lead to refined systems, able to analyse large datasets effectively. We suggest that in the future, combining large-language models with classical AI approaches such as those presented here may help solve more complex scientific and mathematical problems.

²The form is:

$$nQ(n) = \alpha n^{\frac{3}{2}} + \beta n + \gamma n^{\frac{1}{2}} \quad (5)$$

References

- [1] David Brink. A (probably) exact solution to the Birthday Problem. In *The Ramanujan Journal*, 2012.
- [2] Matthew Conroy. Determining quadratic functions. <https://sites.math.washington.edu/~conroy/m120-general/quadraticFunctionAlgebra.pdf>, 2012.
- [3] Miles Cranmer. Interpretable Machine Learning for Science with PySR and SymbolicRegression.jl. In *arXiv:2305.01582*, 2023.
- [4] Miles Cranmer, Alvaro Sanchez-Gonzalez, Peter Battaglia, Rui Xu, Kyle Cranmer, David Spergel, and Shirley Ho. Discovering symbolic models from deep learning with inductive biases. In *Proceedings of the 34th International Conference on Neural Information Processing Systems*, 2020.
- [5] Brian Falkenhainer and Ryszard Michalski. Integrating Quantitative and Qualitative Discovery: The ABACUS System. In *Machine Learning*, 1986.
- [6] Seungwoong Ha and Hawoong Jeong. Unraveling hidden interactions in complex systems with deep learning. In *Scientific Reports*, 2021.
- [7] Bruce Koehn and Jan Zytkow. Experimenting and theorizing in theory formation. In *Proceedings of the International Symposium on Methodologies for Intelligent Systems*, 1986.
- [8] Patrick Langley. BACON: A production system that discovers empirical laws. In *Proceedings of the Fifth International Joint Conference on Artificial Intelligence*, 1977.
- [9] Patrick Langley, Herbert Simon, and Gary Bradshaw. Heuristics for empirical discovery. In *L. Bolc (Ed.), Computational models of learning*, 1987.
- [10] Patrick Langley, Herbert Simon, Gary Bradshaw, and Jan Zytkow. Scientific Discovery: Computational Explorations of the Creative Process. 1987.
- [11] Bernd Nordhausen and Patrick Langley. A robust approach to numeric discovery. In *Proceedings of the Seventh International Conference on Machine Learning*, 1990.
- [12] Takashi Washio and Hiroshi Motoda. Discovering admissible models of complex systems based on scale types and identity constraints. In *Proceedings of the Fifteenth International Joint Conference on Artificial Intelligence*, 1997.

A Appendix / supplemental material

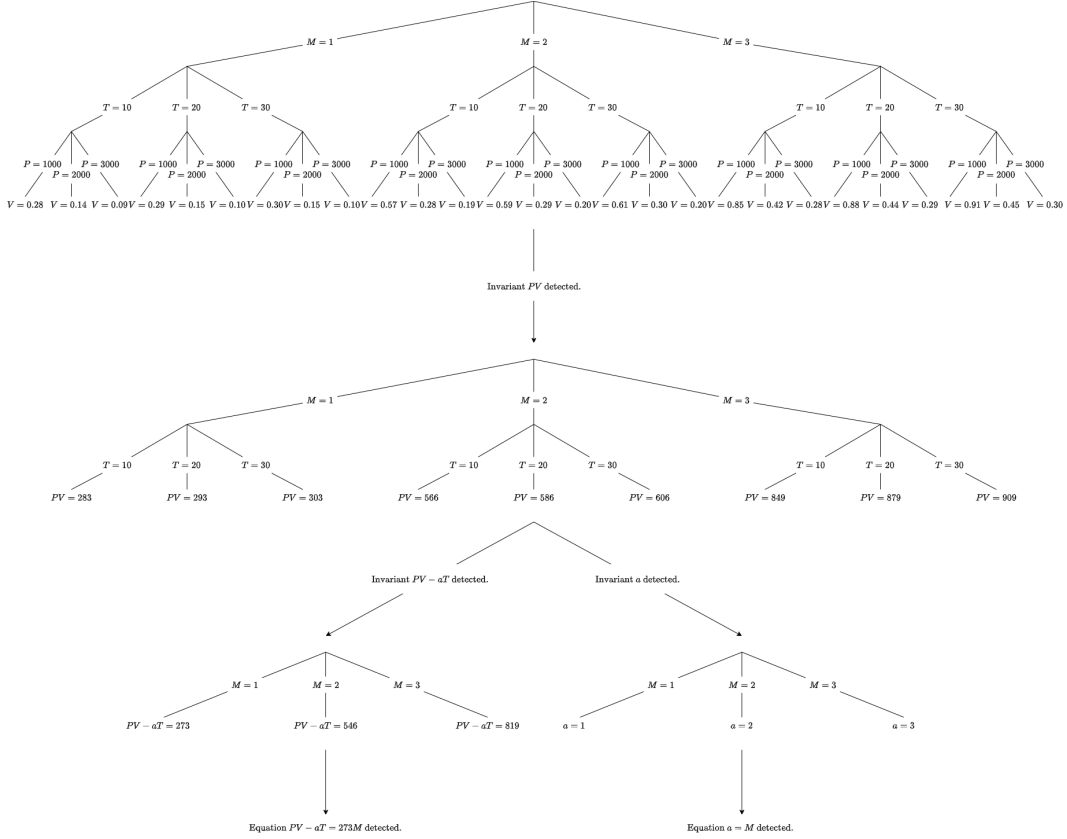


Figure 3: Consider the Ideal Gas Law $V = \frac{M(T+273)}{P}$. It is placed in a noiseless environment where V is a dependent variable, P , M and T are independent variables each consigned to 3 specific values - ultimately giving 27 values for V . At the lowest level it runs the simple BACON.1 heuristic between P and V , here the sets share the same value in M and T . For example, the first set has $T = 10$, $M = 1$ and are the 3 points defined by $P = 1000$, $V = 0.28$, $P = 2000$, $V = 0.14$ and $P = 3000$, $V = 0.09$. All 9 sets in this layer determine invariant PV . BACON.3 detects the agreement, then forms a new tree with PV as the dependent variable. It then proceeds to run a new heuristic check between PV and T . Here all 3 sets determine $PV - aT$ is invariant for new dependent variable a . As both a and $PV - aT$ can be a function of M , BACON.3 forms two new trees at this level. The last BACON.1 heuristic check finds $a = M$, $PV - aT = 273M$. These are collated to form the Ideal Gas Law.

A.1 Design, Algorithmic Details and Implementation

This section will outline the algorithm and mechanisms implemented to reproduce BACON, including adaptations to make it significantly more noise-resilient. Then we will discuss key features of the implementation to emphasise understandability, especially compared to PySR.

B BACON.1

BACON.1 is an algorithm that deciphers the relationship between two symbols X and Y with associated datasets $d(X)$ and $d(Y)$ and no prior knowledge of how they interact. It is an iterative process, where the two symbols X and Y are steadily combined to attain the correct form. This is demonstrated in Table 1 for symbols D and P . We built BACON.1 due to it being the simplistic backbone of BACON.

It consists of two components. The **heuristic-layer** determines the next step in the process from symbols X and Y . It outputs either a function $f(X, Y)$ with associated dataset $f(d(X), d(Y))$, or

null if X and Y don't share a valid relation. The **managing-layer** lies on top of the heuristic-layer and chooses which variables and hyperparameters to pass to it. It controls the complexity and scope of relationships that can be found.

B.1 Heuristic-layer

Overview

The heuristic-layer is handed two symbols, X and Y , respective datasets $d(X)$ and $d(Y)$ and hyperparameters δ , ϵ and c_{val} . The order of its run is in flowchart Figure 4.

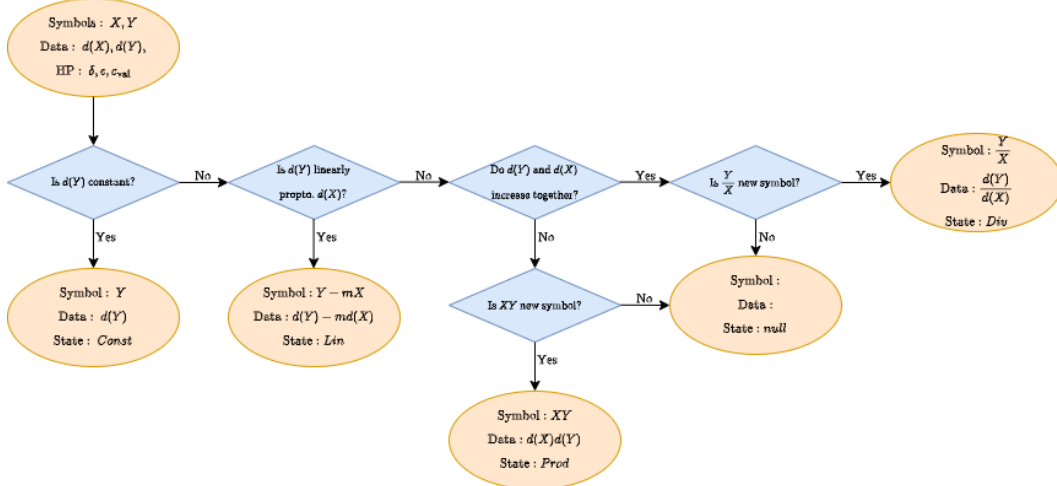


Figure 4: The full algorithm for the heuristic-layer. When the data is returned via $f(d(X), d(Y))$ the function is acted on element-wise. Also, the state return for a linear relationship, includes a list of the parameters of the relationship, such as the gradient and y -intercept. The consistency check determines if $d(Y)$ is constant, and the linear proportionality check determines if the data is related linearly. The novel relation check determines if a new symbols has been found. These checks implicitly use the hyperparameters the heuristic-layer is initialised with, where appropriate.

The logic behind the consistency check, linear proportionality check and novel relation check are described below.

Consistency check

For $M_Y = \text{mean}(d(Y))$. If for each value in $d(Y)$:

$$M_Y(1 - \delta) < \text{value} < M_Y(1 + \delta) \quad (8)$$

then Y is defined as constant. Increasing δ increases the frequency in which relationships are found. To deal with negative values, the absolute value of the datasets are used throughout. A practical example is explained through Figure 5.




Figure 5: For symbol Y with $d(Y) = \{0.80, 0.95, 1.25\}$ so $M_Y = 1$, the consistency rule holds for the green dotted boundaries defined by $\delta = 0.30$, but not for $\delta = 0.15$. If the consistently rule holds, the quantity is defined to be invariant - here in the case of $\delta = 0.30$, Y would be defined as invariant with value of M_Y .

Linear proportionality check

X and Y are compared via the correlation coefficient³. If $|r| \sim 1$ then X and Y are linearly proportional. $r > 0$ demonstrates that the sets increase or decrease together and they should be divided, $r < 0$ displays one increases whilst the other decreases and their product should be taken. Graphically this is shown by Figure 6.




Figure 6: A demonstration of the different values of r for ordered datasets $X = \{x_1, x_2, x_3\}$ and $Y = \{y_1, y_2, y_3\}$. The pink dotted line is the line of best fit through the points. If $|r| \sim 1$, then X and Y are linearly proportional. A positive r implies they increase together, a negative r shows that one increases whilst the other decreases.

A linear relation is deduced if $1 - |r| < \epsilon$. Increasing ϵ increases the frequency linear relationships are ascertained. If a linear relationship is determined, a check is made that the y -intercept is non-zero. This is as $X + mY = 0$ is equivalent to $\frac{X}{Y} = -m$. The former means variables $X + mY$, and m have to be stored by the system whilst the latter only requires $\frac{X}{Y}$ meaning its quicker and simpler to be handled. The threshold is determined by variable c_{var} . For approximated linear relationship $d(Y) = md(X) + c$, if $|\frac{c}{\text{mean}(d(Y))}| < c_{\text{var}}$ then the y -intercept is classified as zero.

Novel relation

The iteration process discovers new forms. To make sure it doesn't recover an already found symbol, the output of the product/division step must be checked. For example, if the heuristic-layer was passed symbols X and XY by the managing-layer, and XY does not pass the consistency or linear proportionality check but has $r > 0$. The new symbol would be $\frac{XY}{X} = Y$. As Y has already been seen this is classified as an invalid symbol, and the heuristic-layer would return state null.

B.2 Managing-layer

A visualisation of the managing-layer and its interaction with the heuristic-layer is in Figure 7. The goal is to uncover the most accurate relation between X and Y , whilst running the heuristic-layer the minimal amount of times. There are two key variables throughout. j is the iteration counter. It controls the complexity of the allowed equations. Variables s_ϵ and s_δ are the scale factor ϵ and δ are multiplied by, if no equation has been found once the iteration counter maxes out or no relationship is found between X and Y .

Order of events

It runs a sequential search where the new relationship - Z in Figure 7 - is compared first to variable X then to Y - even if X finds a relation. This is done to allow X and Y to be tested as a linear relationship with each found variable, whilst not searching exhaustively, in the aim of running the heuristic-layer as little as possible. My method works under an assumption described through the following example. If the true relation to find is $\frac{X^3}{Y}$, and the heuristic-layer run of X and $\frac{X^2}{Y}$ has outputted this symbol. By the order above, the heuristic-layer between $\frac{X^2}{Y}$ and Y now runs. My

³For ordered datasets $d(X) = \{x_i : i \in 1, \dots, n\}$ and $d(Y) = \{y_i : i \in 1, \dots, n\}$ with respective means M_X and M_Y , the correlation coefficient r is defined by:

$$r = \frac{\sum_{i=1}^n (x_i - M_X)(y_i - M_Y)}{\sqrt{\sum_{i=1}^n (x_i - M_X)^2} \sqrt{\sum_{i=1}^n (y_i - M_Y)^2}} \quad (9)$$

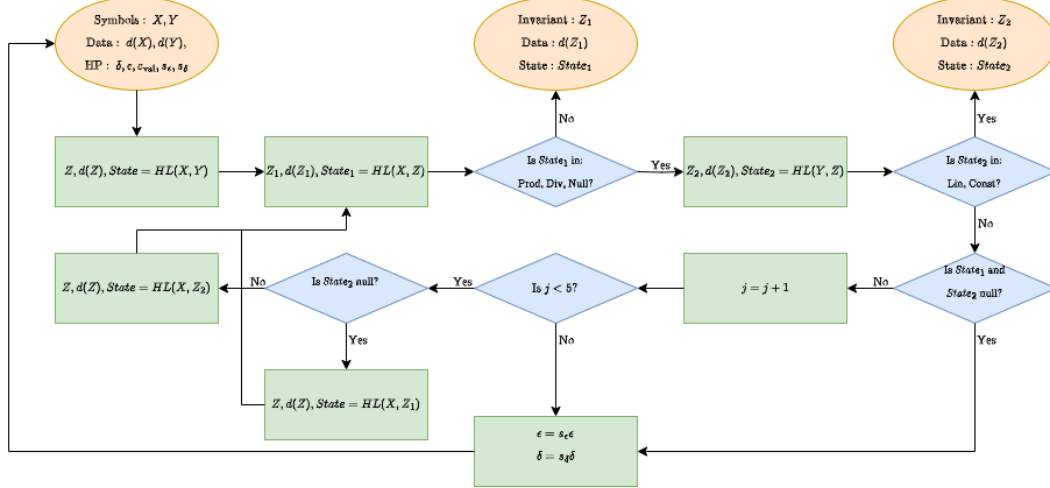


Figure 7: Flowchart describing the managing-layer of the BACON.1 algorithm. The heuristic layer is run whenever HL is called. It implicitly inputs the δ , ϵ and c_{val} hyperparameters as well as the datasets of the symbols being inputted.

assumption is that the relationship found is a product, which is invalid as X^2 is an invalid constant as it would have already be found by X being constant. This allows the next run of the heuristic-layer to occur between $\frac{X^3}{Y}$ and X confirming $\frac{X^3}{Y}$ as the invariant. We are yet to find a counter-example here.

Iteration Counter

j increments every-time the heuristic-layer returns a value. This adds a complexity cap to the relationships found. This is demonstrated in Figure 8. Every new symbol is of the form $X^n Y^m$ for




Figure 8: All 10 possible relations for a BACON.1 run capped at $j = 4$. Note this is the minimum required to find Kepler's Law.

$|n| + |m| \leq j + 1$ and $n, m \in \mathbb{Z}$. Each of these symbols can be in a linear relationship with X and Y .

Scale Factors

The scale factors cause auto-increasing ϵ and δ if the managing-layer concludes without finding an invariant. In noisy environments the hyperparameters for error tolerance needed vary between layers in the tree. Allowing the system to auto-update itself means a relationship will inevitably be found for

every set. These are governed by s_ϵ and s_δ . The updated ϵ and δ get multiplied by this scale factor if the managing-layer is rerun. Tests were done with incremental rather than multiplicative scale factors. On noisy datasets, the ϵ and δ needed at the final layer of the tree, when only 2 variables remain, are experimentally orders of magnitude larger than earlier layers. The multiplicative factor gets to these magnitudes quicker than an incremental approach. Larger s_δ will increase the programs tendency to find product/division relations, likewise for s_ϵ and linear relations.

B.3 Strengths of BACON . 1

BACON . 1 is fast. Consisting of only simple checks between two sets of 3 datapoints allows a run of the managing-layer BACON . 1 to take ~ 0.03 seconds. On noisier datasets, reruns of the managing-layer are more likely. Due to the speed of BACON . 1 multiple reruns can occur without a large time deficit. Moreover, reducing this speed was the reasoning behind certain design decisions such as the complexity counter and the order of events in the managing-layer.

The alternative BACON . 6, discussed in section C, is the other method Langley devised of a similar task and it sacrifices speed for accuracy. BACON . 1 is about $4\times$ faster for a single run whilst not requiring prior knowledge of the relationship between the symbols which BACON . 6 does. It allows BACON . 1 to work in more general spaces.

This combination makes BACON . 1 versatile and adaptable to different datasets no matter the noise. It thrives under modern computing which runs its algorithms quickly.

B.4 Weaknesses of BACON . 1

BACON . 1 needs at least 3 datapoints to infer any relationship. With only two datapoints, multiple relationships can be formed that fit the data equally well. For example if $X = \{2, 3\}$ and $Y = \{4, 9\}$, relationships $Y = X^2$ and $Y = 5X - 6$ are both valid, though a third datapoint would distinguish between these - or give another relation. This tautological nature is an essential limitation of BACON with no fix apart from gathering more data. Furthermore, this reasoning is why BACON can't deal with relations that are sums of three terms ie. $f(X, Y) + ag(X, Y) + bh(X, Y) = \text{constant}$, for functions f, g, h . A quadratic function is uniquely determined by 3 points [2], however terms of higher powers may hold the same relation. For example, points $X = \{1, 2, 3\}$, $Y = \{1, 8, 27\}$ satisfy both $Y = X^3$ and $Y = 6X^2 - 11X + 6$. To distinguish, an additional datapoint is needed.

The managing-layer only compares each term found in Figure 8 with X and Y to save time. This limits the relationships found. For instance, it couldn't find the linear relationship $X^2Y^3 - aXY$. This is changeable for a computational complexity and time trade-off. For the datasets used in this project it is not a necessary change.

B.5 Relation to Langley's BACON . 1

Heuristic-layer

The heuristic-layer is almost identical to Langley's. Both the consistency and linear proportionality tests are in Langley's BACON . 1. However, we add the hyperparameter c_{val} . This was done experimentally when we were testing on noisier data, and we discovered linear relationships were being found spuriously. Consider Ohm's Law (details in subsection F.2), the first layer requires the relationship $IL - aI$ to be found. Instead, $I - aL$ appeared consistently in noisy data with few instances of $IL - aI$. Reducing ϵ stopped the few instances $IL - aI$ appearing which instead formed relation IL . However, increasing c_{val} to 3, stopped $I - aL$ appearing at all whilst not affecting $IL - aI$.

Managing-layer

Langley doesn't explicitly describe his managing-layer, but must have one in order for his heuristic-layer to function. He also makes no mention to an iteration counter or the scale increases for δ and ϵ . His program only terminated with an explicit stop command, or when a relationship was found. When he dealt with noiseless data this wasn't an issue as a relationship is almost always found. Our method means we can successfully deal with noisy data, which needs different hyperparameters at different layers in the tree. This was the most crucial upgrade to BACON . 1 to allow it to perform well on noisy data.

C BACON.6

BACON.6 was the last iteration in the BACON family. It was described by Langley as a “hill-climbing method for dealing with noise” [9]. BACON.6 was implemented due to Langley’s belief in its noise-handling. In addition, it is the only understandable classical algorithm we’ve found that is adaptable to find invariants that contain functions that are trigonometric or exponential.

The difference with BACON.1 is that BACON.6 requires prior knowledge about the type of relationship the two symbols share. The algorithm is smaller than BACON.1 due to this prior knowledge.

C.1 Method

Given symbols X and Y with datasets $d(X)$ and $d(Y)$ and a basic understanding of their relationship, for example⁴, $Y = \alpha X^2 + \beta X + \gamma$. BACON.6 iterates through a set of possible values for α and β (γ – as a constant – is ignored at this stage as it is combined with another constant after the algorithm is done), saving those with the greatest correlation coefficient. The hyperparameters are an initial N which can be any positive number, an $n_{\text{threshold}}$ which is set to 2, and a p which determines how many iterations the process performs. It is as follows:

1. Initialise 9 instances of $\alpha X^2 + \beta X$ with α and β taking values of all 9 possible pairs $\in \{-N, 0, N\}$. These are $(-NX^2 - NX)$, $(-NX^2)$, $(-NX^2 + NX)$, ...
2. Calculate the correlation coefficient of $d(Y)$ and all 9 combinations above with $\alpha d(X)^2 + \beta d(X)$. Retain the α and β corresponding to the greatest $n_{\text{threshold}}$ of them. This measures which combinations fits the data best.
3. Initialise 9 instances of $\alpha X^2 + \beta X$ with α and β taking values of all 9 possible pairs $\in \{-\frac{N}{2}, 0, \frac{N}{2}\}$. Add the retained α and β to each pair, creating 18 combinations each an addition of a retained expression and a new combination.
4. Calculate the correlation coefficient of $d(Y)$ and all of these 18 combinations as before. Retain the best $n_{\text{threshold}}$ of them.
5. Step (3) and (4) are repeated, each time halving the N that the new combinations instantiate from. The program terminates when $N < \frac{1}{2^p}$. The output is the α and β which gives the greatest correlation coefficient.

We now have a linear relation between Y and $\alpha X^2 + \beta X$. `polyfit` from NumPy is used to perform linear regression which gives the optimal gradient m and y -intercept c to finalise the exact relationship. Including additional variable γ from earlier, would increase the algorithm time unnecessarily as it still would have combined with c at this stage.

Langley never incorporated BACON.6 in the wider BACON search tree with BACON.3 or BACON.5. He only detailed a small number of instances of BACON.6 solving examples like this. His main example was the less convoluted $Y = 3X^2 + 2X + 1$. Hence, the fact our build of BACON.6 can solve the comparatively complicated birthday problem⁵ displays a faithful reconstruction of Langley’s work. Moreover, we have implemented BACON.6 into the search tree via using a general form for BACON.6 to solve on - something Langley never did. For coefficients α, β, γ and symbols X, Y the form is:

$$Y = \alpha X + \frac{\beta}{\gamma + X} \quad (10)$$

However, it ultimately wasn’t efficient for reasons detailed in subsection C.2.

C.2 Weaknesses of BACON.6

Adding more variables, and increasing the generality of the form comes at an expensive time and computational complexity trade-off. Recall in the method 9 combinations of α and β are generated each step. This extrapolates to 3^n combinations being formed for n the number of variables.

⁴The RHS of the example can contain any type of function such as trigonometric or exponential or even those that contain Y . It is set as a polynomial here for simplicity.

⁵We did confirm this reconstruction could solve $Y = 3X^2 + 2X + 1$. However, as this is also solved by PySR it doesn’t demonstrate Classical AI outperforming a DNN like the birthday problem does.

Completing the correlation calculations then takes exponentially longer with each additional variable, as there is $\mathcal{O}(3^n)$ computational complexity. 3 variables was found to be the maximum for BACON . 6 to solve in good-time. Even then it took $4\times$ longer than BACON . 1. This makes BACON too restrictive in what it can solve. For example, it couldn't solve Ohm's Law as it involves a squared term for D which corresponds to a square for Y in equation 10.

The reason for having to use a generalised form is symptomatic of the true issue. Having to input a predicted form of the equation goes against the premise of BACON inferring everything from the data. The user often will have no knowledge of the data they're using rather just a dataset.

Another issue is that Langley's claim that BACON . 6 is especially adept at handling noise doesn't seem true. For example, even at 0.5% relative noise, the found equation from the birthday problem distorts to:

$$Q_{\text{BACON.6}} = 0.968\sqrt{n} + 1.307 - \frac{2.904}{\sqrt{n}} + \frac{1.732}{n} \quad (11)$$

A form dissimilar to equation 6, with an MSE of 0.002 to the noiseless Q - 7 orders of magnitude worse. Results shown in subsection F.6 tended to only get a single order of magnitude worse with every 0.5% noise added to the dataset. This is then anomalous in how poor the performance is. We suspect this is so because BACON . 6 overfits to the noisy data through the use of α and β coefficients being too precise.

In all, BACON . 6 is slower and limited in the forms it can find compared to BACON . 1. It has strengths in noiseless environments but can't adapt to noise well enough to be used in our main model BACON . 7 (in subsection E). Langley's other methods didn't handle noise well, so with perspective to those - and the fact a results is always given (via the coefficients), this may have given him cause to praise BACON . 6's error tolerance.

C.3 Relation to Langley's BACON . 6

The program was implemented exactly as described in [9] with the addition of the parameter p . Langley used a fixed threshold to stop the iterations.

D The Space of Data

Langley designed both BACON . 3 and BACON . 5 as ways to traverse a multivariable dataset. Both were flawed when dealing with noisy datasets. As such, we've designed our own novel methodology to explore the space. It uses BACON . 3's idea of traversing layer-by-layer, and BACON . 5's idea of symmetry.

This section will describe how we have approached this problem, our solution and comparisons with Langley's versions as well as why they are flawed. It also discusses why BACON . 5 may be the best method on large datasets.

D.1 Method

At each layer in the tree, a 3-step process is followed:

1. Check if the layer is sufficiently constant. If it is, terminate the tree with the symbol representing the layer. At the last layer, this check is not run.
2. If it is not sufficiently constant, calculate the expression for each set in the layer by inputting the set into BACON . 1 or BACON . 6. Use a layer-method to determine which expressions is chosen.
3. Average the layer to form a new layer. This is crucial in BACON's ability to determine the true equation in noisy environments as displayed in subsection F.6.

Sufficiently constant

Similar to δ earlier, parameter Δ is introduced to prune a tree if the bottom layer is sufficiently constant. If the bottom layer has symbol Y and datapoints $d(Y)$. For $M_Y = \text{mean}(d(Y))$, if $(1 - \Delta)$

proportion of the values⁶ in $d(Y)$ satisfy

$$M_Y(1 - \Delta) < \text{value} < M_Y(1 + \Delta) \quad (12)$$

then the invariant is determined as $Y = M_Y$. The comparison between Y and other variables is skipped. Throughout the remainder of the project $\Delta = 0.01$ but is user-configurable. It saves time. For example, in Ohm's Law, the invariant $IL - aI$ is discovered in the first layer. A new tree with a uses this test to determine $a = 3$ without comparing to other variables T and D .

Layer-methods

We have devised various methods to pick an expression from the list of possible expressions. It was one of the main novel contributions to improve BACON. They are listed in Table 2 with their name, an explanation and the time taken to operate on the first layer of Black's Law with 0.5% noise. This example has 81 datapoints and 27 sets leading to 2 expressions which have to be ranked.

Method	Time	Description
bacon.3	3×10^{-6} s	Requires the sets to all find the same expression. If this does not happen the program terminates.
popular	7×10^{-6} s	Selects the expression which appears most frequently - if there is a tie, defer to min_mse.
min_mse	0.37s	Apply an expression to a set to gain 3 datapoints. Normalise the datapoints. Take the mean squared error (MSE) between the normalised datapoints and their mean ⁷ . Sum over each set in the layer to get a total MSE for the expression. Do this for all expressions, the expression with the lowest total MSE fits the dataset best and is selected.
gp_ranking	0.85s	Apply an expression to all the sets. Gaussian Process (GP) Regression is performed on this new dataset to find the optimal GP that fits the data. The GP is sampled, and the signal-noise-ratio of the sample calculated. This done for each expression. The highest score implies the least noise, and best fit of the expression to the data. This expression is selected.
user_input	N/A	The user selects the expression they'd like to use from those found by the sets. Used for debugging.

Table 2: A table displaying the custom layer-methods we devised as well as an explanation of how they select an expression. The time they each take to rank a layer with 81 datapoints and 2 found expressions is also given.

bacon.3 and popular are fastest as a dataset doesn't have to be reconstructed with the found expression applied to the last layer. This has to happen for min_mse and gp_ranking. min_mse takes less than half the time as gp_ranking due to the complexity in performing GP regression. They also gave near identical results in choosing the correct expression. It displays that black-box techniques aren't necessarily better than Classical techniques especially on simple tasks like calculating MSEs. popular and min_mse were used throughout the results in subsection F.6.

Averaging

From Figure 3 we see the tree gets smaller at each iteration of the tree. Once an expression has been determined via the layer-method it is applied to each set. Then in the set, there are 3 datapoints from the expression which are averaged. This becomes a new datapoint, reducing the overall number of datapoints by 3. For example, the first set in the first layer in Figure 3, has $T = 10$, $M = 1$ and are the 3 points defined by $(P = 1000, V = 0.28)$, $(P = 2000, V = 0.14)$ and $(P = 3000, V = 0.09)$. PV is determined and the datapoints found are $PV = 280, 280, 270$ ⁸, this would be averaged to

⁶Eg. if $\Delta = 0.02$, then 98% of the values in $d(Y)$ must satisfy the Δ relationship above. This allows it to deal with extremely noisy points due to variance in sampling noise from a Gaussian distribution. Without this, the test rarely passes due to these outliers.

⁷For found variable X with dataset $\{x_i : i \in 1, \dots, n\}$ and mean M_X , the MSE is defined as:

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (x_i - M_X)^2 \quad (13)$$

⁸In Figure 3 the dataset is noiseless, and so with the full decimal expansion of V , PV would be 283. However, this is a demonstration of what happens when there's noise.

$PV = 277$. Then the new tree which has PV at the bottom layer would have $PV = 277$ for the datapoint with $T = 10, M = 1$. We hypothesise, that this averaging places a crucial role in BACON stripping away noise. This is elaborated on in subsection E.

D.2 Why are multiple methods needed?

Consider the Ideal Gas Law at 3% relative noise with the hyperparameters found in `args/ideal/ideal30.json`. A run through the tree with the popular layer selected looks like:

1. In the first layer the sufficiently constant check is not passed by V . Instead expressions are found, and the 9 sets determine: PV invariant 7 times and $PV^2 - aV$ and $V - aP$ are both found once. PV is carried forwards by popular layer and applied universally to the dataframe. The sets are averaged reducing the dataframe size.
2. The sufficiently constant check is not passed by PV . Then, between PV and T , the 3 sets find the relation PV , $PV - aT$ and $\frac{T}{PV} - aT$. As these all appear once, we defer to the MSE calculations. These are respectively 0.0043, 0.0027, 0.0032 causing $PV - aT$ to be used. Thus this alternative MSE mechanism is needed to decide. The sets then average once again.
3. As the tree is on the final layer, the sufficiently constant test is not run. Between $PV - aT$ and M the single relationship found is $\frac{M}{PV - aT} = 0.0036$. Between a and M the relation is $\frac{M}{a} = 1.04$, altogether yielding $V = \frac{M(0.96T + 273.57)}{P}$ which best fits this noisy data.

In general `min_mse` works best with little amounts of noise and as a tiebreaker for `popular`. Once the data gets noisier `min_mse` ends up getting tricked. In the example above at the first layer, the MSE for PV is 0.006, $PV^2 - aV$ is 0.004 and $V - aP$ is 0.036. $PV^2 - aV$ would have got picked, which would stop the true equation being formed. This happens as $PV^2 - aV = \text{const} \implies PV = a + \frac{\text{const}}{V}$. For V significantly larger than const , the RHS is approximately constant. This is another demonstration of the need for c_{val} from subsection B.1. Increasing it to 0.1 from 0.02 stops $PV^2 - aV$ becoming a valid invariant and thus retains the ability to use the `min_mse` method.

D.3 Weaknesses of layer-methods

Both `popular` and `min_mse` have their flaws. Experimentally, it is easy (on multiple repeats) to find a set of hyperparameters where the real equation can be determined if the correct invariants are chosen at each stage (done by `user_input`). However, once the data becomes noisier, it is often the case it is neither the most popular equation, nor satisfying `min_mse`. In addition, with this noise, different layers are solvable by different layer-method which so far has been too difficult to implement in a generalised manner.

The other issue is that the order of the layers matter. Consider Black's law again:

$$T_f = \frac{M_1 T_1}{M_1 + M_2} + \frac{M_2 T_2}{M_1 + M_2} \quad (14)$$

Initially between T_f and T_2 there is invariant $T_f - aT_2$. If instead the first layer was between T_f and M_1 the equation rearranges to $T_f M_1 + T_f M_2 - M_1 T_1 = M_2 T_2$, leading to invariant $T_f M_1 - aM_1 - bT_f$ which isn't solvable for reasons explained in subsection B.4. However, a brute-force approach could be generated to keep varying the input order until a form that is solvable is entered. This is a potential extension to BACON, and similar to what `Fahrenheit` – mentioned later in subsection H.1 – performed.

D.4 Relation to Langley's BACON.3 and BACON.5

Langley's BACON.3 required ubiquity along the expressions found by the sets. One of our layer-methods works along this premise, named `bacon.3`. With noise, the sets don't always agree, and with more noise it becomes more likely they don't agree. This limits its use.

BACON.5 introduced symmetry. At each level only a single set is needed to form an expression which can then be applied to the sets in the layer by symmetry as they all share the same relation. This directly inspired me to gain all the expressions, with the best (determined by the layer-method) applied to the rest of the sets in the layer.

Langley’s method required picking the datapoints that are used to produce the sets before any calculations were made. There are many ways to do this. For example, Figure 9 demonstrates possible choices for a dataset with 3 variables and 9 datapoints assuming that the set is the first 3 datapoints from the bottom left. There are 3 possible configurations for this initial set, giving a total of 27 possible index choices.




Figure 9: Possible ways to pick datapoints in BACON .5 when the set needed is in the bottom left. The chosen datapoints are in black. As the initial set could also be from the centre or right branch, there’s a total of 27 possible ways to choose.

Extrapolating for an $n + 1$ variable system, there are $3^{n(n-2)}$ ways to do this⁹. This is an exponentially increasing choice. If made correctly, it improves the likelihood of BACON .5 leading to the correct expressions. This method also reduces the datapoints needed. Now the program only needs to consider $3 + 2(n - 2) = 2n - 1$ datapoints. This is significant in very large datasets as it requires small amounts of perfect data. BACON also performs better in smaller datasets and being able to create this environment is desired. If an optimal method could be discovered for picking branches, BACON .5 would be effective on large datasets.

We could not find a method for this that works quickly and efficiently. Langley made this in his noiseless environment where datapoint-picking was arbitrary. With noise this is not the case. Multiple datapoints are significantly worse due to random Gaussian noise that picking them at any stage is tantamount to failure.

We used a Gaussian Process (GP). We rank each datapoint by fitting a GP to the dataset with that datapoint removed, then calculating the signal-noise-ratio of a sample from that GP. The datapoints corresponding to the GPs with the largest signal-to-noise ratio are noisiest, and should not be picked. Doing this for all 3^n datapoints is time-consuming, computationally complex and goes against the purely heuristic-principles of BACON. However, this does show a situation where Classical AI and modern black-box techniques can be combined.

E BACON .7

E.1 Overview

BACON .7 is the name we have given for our version of BACON. It uses our layer-method design as described in section D. It uses BACON .1 as the Space of Laws.

At the Space of Laws level, BACON .6 with a general equation form and an implementation of PySR were tested. The latter struggled consistently as 3 datapoints is not enough for it to function efficiently - a feature seen repeatably through the results in subsection F.6. BACON .6 tends to overfit on noise, as well as taking too long when used with a general form, both discussed in subsection C.2. However, there is a question: if BACON .1 can generate a relationship, can BACON .6 improve the found constants in said relationship? This is tested in subsection E.2.

The project was coded with an eye on understandability and clarity. The system had to be straightforward and reproducible. Details on how this was done are given in Section E.4.

⁹Extrapolating for an $n + 1$ variable system, with n independent variables, 1 dependent variable, and calling the total choices $c(n)$. From the first independent variable down, there are 3 possible values of this variable that the initial set can take. In the branches of this variable coinciding with the 2 values not chosen there is 3^{n-2} possible options as any index can be picked. The equation then becomes $c(n) = 3 \times 3^{n-2} \times 3^{n-2} \times c(n - 1)$. Noting $c(2) = 1$ and iterating this process, $c(n) = 3^{n-2+2(n-2+\dots+1)} = 3^{n(n-2)}$.

E.2 Combining BACON . 1 and BACON . 6

This is tested on the Ideal Gas Law. On the conclusion of the tree-search, and as elaborated on in subsection F.1, the equations found are:

$$\frac{M}{a} = 1, \quad \frac{M}{PV - aT} = \frac{1}{273} \quad (15)$$

Reworking the RHS constants as variables α and β respectively, the following is gained:

$$V = \frac{M}{P} \left(\frac{T}{\alpha} + \frac{1}{\beta} \right) \quad (16)$$

This is the ideal form to input into BACON . 6. Evaluating on the Ideal Gas Law with 3% noise without BACON . 6 warrants the form

$$V_{\text{BACON.7}} = \frac{M}{P} (0.96T + 273.57) \quad (17)$$

with MSE of 1.6×10^{-4} to $V + n$ (the noisy data, detailed in section F.4). Comparatively, BACON . 6 integration gives form

$$V_{\text{BACON.7}} = \frac{M}{P} (1.02T + 280.18) \quad (18)$$

with MSE of 1.4×10^{-4} to $V + n$. When comparing instead with noiseless V , the MSE difference now becomes 3.4×10^{-7} using BACON . 1 but 1.2×10^{-4} with BACON . 6. Hence, BACON . 1 is four orders of magnitude more accurate without BACON . 6. As this is a measure against the objective true function, BACON . 1 is unequivocally better at stripping noise from the equation, whereas BACON . 6 overfits on the noisy data. This is a trend similarly seen with PySR in the results, subsection F.6.

E.3 Weaknesses of BACON . 7

BACON . 7 is unable to work at variables that scale orders of magnitude apart. For example, the SIR model with variables that vary at scales of 10^{-8} , 1 and 10^6 could not be used in BACON. This is because it causes the hyperparameters inputted into BACON . 1 to vary between layers preventing the correct expressions to be found. For instance, the first layer may need $c_{\text{val}} = 1$, but the next needs $c_{\text{val}} = 1 \times 10^{-6}$. A future improvement to BACON involves devising a mechanism to vary hyperparameters between layers smartly.

E.4 Implementation

Making the program run on a system had to be straightforward and reproducible, whilst an option for a verbose mode had to be consistently available. The former was accomplished via the ability to write argument files dealing with all variables. An example file, looks like figure 10:

The command to run on a dataset - such as the Ideal Gas Law - with 3% noise is:

```
python3 main.py --dataset ideal --noise 0.03 --args args/ideal/ideal30.json
```

The output of the above - our recommendations for the minimalist approach to understand the order of BACON's implementation - can be seen in Figure 11.

Comparatively, the output to a run of PySR - for the same dataset and noise - is shown in Figure 12. It demonstrates a list of possible equations PySR builds on to present its final output. There's no explanation of how it accomplishes this. In addition, the outputs are not repeatable further reducing understandability.

F Datasets and Results

The datasets used for this evaluation are synthetically generated from three well-known equations described below. In each equation set, the values for the independent variables were based on the values Langley displayed in his book [10].


```

1 // args/ideal/ideal30.json
2
3 {
4   "layer_method": "min_mse",
5   "layer_args": {"verbose": true},
6
7   "laws_method": "bacon.1",
8   "laws_args": {"epsilon": 0.025,
9                 "delta": 0.04,
10                "c_val": 0.5,
11                "epsilon_scale": 1.05,
12                "delta_scale": 1.2,
13                "verbose": false},
14
15   "data_space_args": {"Delta": 0.01,
16                      "verbose": false}
17 }

```

Figure 10: Typical argument file for this project. This can be found in the codebase at args/ideal/ideal30.json, as it was used for the Ideal Gas Law with 3% noise. The simplistic nature of this implementation is key for understandability and reproducibility.

```

Ranking layer: Expressions found with associated popularity are:
                {P*V: 9}
Ranking layer: Proceeding with P*V
Ranking layer: Expressions found with associated popularity are:
                {P*V: 2, P*V - T*a: 1}
Ranking layer: Iteratively ranking the found expressions:
                P*V has average mse 0.004335870878343652
                P*V - T*a has average mse 0.002732361384888301
Ranking layer: Proceeding with P*V - T*a
~~~~~

The constant equations found are:
1.03911145217940 = M/a
0.00365531074561991 = M/(P*V - T*a)
~~~~~

Final form is V = 273.574552094724*M*(0.00351772732169719*T + 1.0)/P
with loss 0.00015958807925621198.
~~~~~

Program took 1.65s!

```

Figure 11: A potential output when running an instance of our BACON. This gives the best trade-off of demonstrating the process of BACON whilst remaining succinct. More details can be outputted (via changing the args file) to detail the logic behind each instance of the Space of Laws, and the overall Space of Data - rather than just each layer.

E.1 The Ideal Gas Law

The Ideal Gas Law is:

$$V = \frac{M(T + 273)}{P} \quad (19)$$

V is volume, P is pressure, M is moles and T is temperature (in this basis the gas content R is 1). M is $\{1, 2, 3\}$, T is $\{10, 20, 30\}$ and P is $\{1000, 2000, 3000\}$. This causes V to vary between 0.09 and 0.91. The ideal gas law has a linear relationship in the numerator alongside a couple of instances of product/division meaning its a useful baseline test. When approached via consecutively introducing variables $V \rightarrow P \rightarrow T \rightarrow M$ (forming the tree in Figure 3) noiselessly, two equations come out of the system:

$$\frac{M}{a} = 1, \quad \frac{M}{PV - aT} = \frac{1}{273} \quad (20)$$

These are combined via dummy variable a to form the Ideal Gas Law.

Hall of Fame:

Complexity	Loss	Score	Equation
1	5.064e-01	1.594e+01	$y = -0.31842$
3	2.523e-02	1.500e+00	$y = 581.48/P$
5	2.145e-04	2.384e+00	$y = M*(290.29/P)$
7	4.902e-05	7.381e-01	$y = ((270.29 + T)/P)*M$
9	4.702e-05	2.087e-02	$y = ((307.99 - (289.71/T))/P)*M$
11	4.645e-05	6.064e-03	$y = ((270.29 + ((T/0.88717) - M))/P)*M$
13	4.600e-05	4.866e-03	$y = (((269 + T) + ((T*0.18127)/M))/P)*M$
15	4.530e-05	7.727e-03	$y = M*((((272.25 - ((M*(3.569/T))*3.5295)) + T)/P)$

Figure 12: The output of PySR. It simply lists the newly found equation at each level of complexity and its associated score and loss with no justification.

F.2 Ohm's Law

Ohm's law is typically seen as $I = \frac{V}{R}$, for V voltage, I current and R internal resistance. An expanded form when considering the law applied to a bar of temperature T , diameter D and length L is:

$$I = \frac{TD^2}{2(L+3)} \quad (21)$$

T is $\{100, 120, 140\}$, D is $\{0.01, 0.02, 0.03\}$, L is $\{0.5, 1, 1.5\}$ whilst 2 and 3 represent the fixed voltage and resistance in the wire respectively. I varies between 0.001 and 0.018. The law contain linear relation in the denominator as well as a square through D . It thus tests many facets of the BACON.1 equation finding system. The system is approached via $I \rightarrow V \rightarrow D \rightarrow T$ yielding equations in a noiseless environment of

$$a = -3, \quad \frac{TD^2}{I(L-a)} = 2 \quad (22)$$

combining along a .

F.3 Black's Law

Black's law as seen is:

$$T_f = \frac{M_1 T_1}{M_1 + M_2} + \frac{M_2 T_2}{M_1 + M_2} \quad (23)$$

T_1 and T_2 are temperatures both taking values $\{50, 60, 70\}$ whilst M_1 and M_2 are masses taking values $\{1, 2, 3\}$. T_f is also a temperature varying between 50 and 70. This is the most involved system solvable by Langley's BACON with 3 linear relationships inside it when approached via $T_f \rightarrow T_2 \rightarrow T_1 \rightarrow M_2 \rightarrow M_1$. The equations found in a noiseless setup are:

$$b = 1, \quad \frac{M_1}{\frac{M_2}{a} - bM_2} = 1, \quad cM_1 = 1, \quad -cM_2 - \frac{T_1}{T_f - aT_2} = 1 \quad (24)$$

which are combined along variables a, b, c .

F.4 Noise

The experiments were run with 0.5% increments in the noise on the dependent variable until 4%. This means for dependent variable D with values $\{d_1, d_2, \dots, d_n\}$, the Gaussian noise added to the scalar d_i is $\epsilon_i \sim \mathcal{N}(0, (0.04|d_i|)^2)$. This new noisy variable is denoted $D + n$.

For reproducibility, the noise added to the dataset is using a NumPy random seed. For repeatability all the argument files used are saved, with the location prescribed to the name of the law. Eg. the Ideal Gas Law run under 0.5% noise has its argument file stored in args/ideal/ideal05json.

F.5 BACON .7 tests

Langley’s approach to testing prescribes that BACON can find an invariant if one combination of the hyperparameters inputted into the system allows the detection of the invariant. We use this approach throughout this section. A similar metric is used for determining if PySR can find the equation - the correct form has to appear in its list of found equations. For reference, the correct form follows the correct placements of the variables whilst the constants and coefficients may vary. Hence for α, β, γ constants the acceptable forms of the previous laws formed by BACON are:

$$\text{BACON}(V) = \frac{M(\alpha T + \beta)}{P}, \quad \text{BACON}(I) = \frac{\alpha T D^2}{(L + \beta)}, \quad \text{BACON}(T_f) = \frac{\alpha M_1 T_1 + \beta M_2 T_2}{\gamma M_1 + M_2}$$

With likewise for the functions formed by PySR. If noisy data for the dependent variable is used to form the predictions, it is defined as $\text{BACON}(\tilde{V})$ for the Ideal Gas Law and similarly for the other two. When calculating the MSE from the predictions of BACON and PySR with the real data, the following notation is used for the Ideal Gas Law: $\text{MSE}(\text{BACON}(\tilde{V}), \tilde{V})$. This displays the MSE between the predicted values of \tilde{V} from noisy data inputted to BACON - effectively how well the model fits the training data. Correspondingly, $\text{MSE}(\text{BACON}(\tilde{V}), V)$ is the difference between the predicted form from BACON fed the noisy data and the actual noiseless V - a measure of how close BACON gets at stripping the noise away and determining the real equation. The same notation is used for PySR.

F.5.1 Setup

Each approach uses BACON .7 discussed in section E. The space of data was chosen between which approach worked best out of `popular` and `min_mse`. As `min_mse` and `gp_ranking` showed similar results previously, the latter was disregarded. The majority of approaches had `popular` as the MSE tests often fails in very noisy ($\geq 3\%$) environment. `popular` also benefits from no calculations done on the data itself saving time.

The experiments were run with 0.5% increments in the noise on the dependent variable until 4% or the true relationship could not be found. For repeatability all the argument files used are saved, with the location prescribed to the name of the law. For instance the Ideal Gas Law run under 0.5% noise has its argument file stored in `args/ideal/ideal05json`.

F.6 Additional Results

Black’s Law

Black’s law varies on a non-log scale. There is exponential-like growth in the MSEs in Figure 13 indicating it is growing at a similar pace to before. *No results could be found to make BACON .7 work at 2% noise or above. Comparatively, PySR finds the exact form of the equation at every noise increment.* This seems to be the optimal environment for PySR; the form of the equation does help its cause. Each coefficient being 1 aligns with the simplistic aims of its final results as discussed in section H.2.

F.7 Discussion

On the Ideal Gas Law

Analysing the equations returned allows a better of understanding of how BACON .7 outperforms PySR. For instance, at 2.5% noise, the form of BACON .7 comes from the two relations in subsection F.1, the first discovers $\frac{M}{a} = 1.024$ and $\frac{M}{PV-aT} = 0.00366$. The constants come from averaging the values found in each step as discussed in subsection D.1. This averaging eliminates much of the noise. Overall displaying final form $V = \frac{M(0.976T+273.488)}{P}$. Comparatively PySR only finds $V = \frac{M(T+270.74)}{P}$. This is because its simplistic aims force the T coefficient 1 - though this is a benefit in experiments on Black’s Law. It is not clear why the other coefficient is 270.74. The model gives no explanation.

On Ohm’s Law

We do not know why PySR cannot find the correct form. For reference at 2.5% noise it finds $I_{\text{PySR}} = D^2(-L + 0.127T + \frac{0.905}{L})$. As the correct form is not found for any of the noise percentages, it indicates a systematic problem rather than a fluke. Instead of reducing noise, it stays in the model

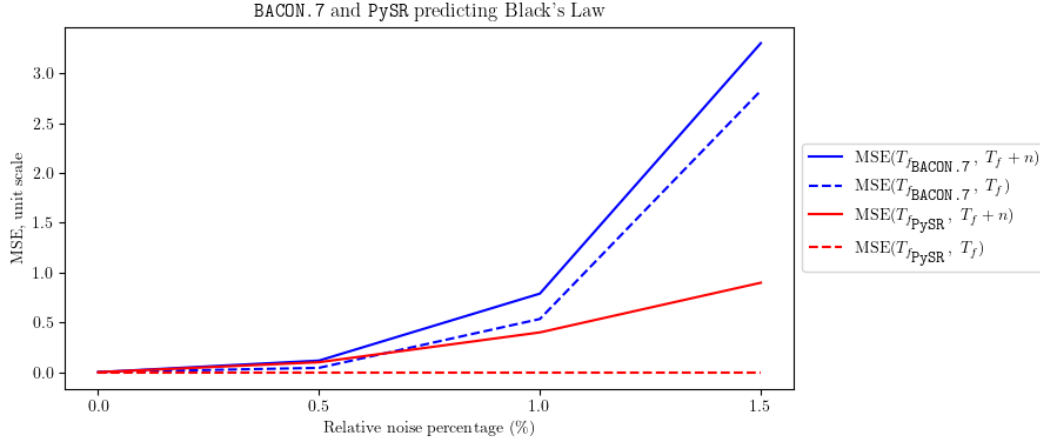


Figure 13: The MSE differences when the above tests are run on Black's Law. PySR outperforms BACON .7. In each case it infers the exact form of Black's Law using its simplistic biases as described in subsection H.2. Additional testing shows it finding this up to 4% relative noise. *This demonstrates the superiority of neural network models, here PySR, when fed enough data.* BACON .7 can't solve greater than 1.5% due to increased noise in the dataset. With 81 values having noise applied, there's a higher likelihood of excessive noise (by random Gaussian sampling) in the datapoints. Even with averaging the noise appears in the last layer, where no combination of the hyperparameters can attain the correct relationship without contradicting previously found expressions in earlier layers.

and this is only known as we know the true form of I . This demonstrates the problem with picking the wrong equation in CSD. Also, the lack of explainability for PySR means we do not know what settings to tweak to fix this issue.

This environment demonstrates that BACON .7 has utility when compared to PySR at this scale. This and the Ideal Gas Law, are on smaller datasets. These are optimum conditions for BACON .7, whilst not giving enough datapoints for PySR to perform effectively. At this scale BACON .7 both strips away noise, and infers the true equation form better than PySR.

On Black's Law

BACON fails as there's a conflict in the hyperparameters to get the correct equations at the final layer in the tree. At 2% the final layer of one of the trees finds $1.74 = 0.296M_1 - cM_2 - \frac{T_1}{T_f - aT_2}$ instead of $1 = -cM_2 - \frac{T_1}{T_f - aT_2}$. To punish this excess linear relationship, ϵ needs to be reduced. However, all attempts there cause the linear relationship for c at the layer above not to form. This contradiction means there's no way of pulling the correct form out of BACON .7.

A reason for this is at previous steps you can choose between multiple expressions as there are multiple sets. As there is only one at the final layer the ability is lost and takes the only expression displayed. A potential improvement of BACON involves more variability at the final step. This is something that we input into MCTS in section G.

Comparison to Langley's BACON

Langley's BACON – which uses `bacon .3`¹⁰ as a layer-method as well as `bacon .1` – could not solve any noise on any dataset, bar 0.5% on the Ideal Gas Law¹¹. As Langley's BACON worked on 0% noise it demonstrated that they had been reconstructed correctly. *Thus, our BACON .7 has been successful in becoming more noise-resilient than Langley's BACON.*

Concluding Comments

For real-life applications the true equation of a dataset will not be known. It is then imperative to be able to trust the equation generated by the model. *For smaller datasets, these results have conclusively*

¹⁰Whilst it was possible to test against BACON .5, Langley's never specified a mechanism for picking initial datapoints. My reconstruction does this randomly and is not reproducible. Thus it was not chosen as the baseline for Langley's BACON.

¹¹Saved under `args/ideal/idealbacon.json`. Likewise the best attempts for Ohm's and Black's law are saved in their respective folders. These only succeed on 0% noise.

shown BACON.7 is more trustworthy than PySR as it is more accurate due to its superior ability to strip away noise. PySR is best on larger, noisier datasets. In these environments, it outperforms BACON.7.

A situation has been found where classical methods outperform modern techniques. It amplifies the need to reproduce, study and understand these seemingly anachronistic mechanisms and see what lessons can be learnt. Also of note, PySR is a complete package, built in collaboration with many engineers over multiple years with well-written documentation and open-source support. BACON.7 was individually built in 7 months with only a 1980s' textbook to guide. With further development, it is possible more use-cases will be found where BACON.7 is able to compete against the state-of-the-art.

An additional note on explainability. we only managed to find the hyperparameters to run the tests up to 4% noise after much trial and error. We could read the output to BACON.7, understand at what stage problems occurred (such as making a linear relationship when it should have been a product) then fix the appropriate hyperparameters. *This demonstrates how explainable models help aid in development.*

G Monte Carlo Tree Search

Multiple layer-methods were created as no method consistently chooses the correct expression when data is noisy. When using the `user_input` method, it seemed each expression did exist, at least until the final layer. Here, we saw from Black's Law in subsection F.6 that the ϵ and δ needed varied from earlier layers. This section uses MCTS to optimise the search through the possible expressions and solve both these issues.

G.1 Method

To use MCTS we need the concept of a node. In chess, this is represented as the position of the pieces on a board at a given go. For BACON we represent it by the expressions found at a given layer. To move to the next node, the legal actions must be determined. In chess, these are the set of legal moves. In BACON these are the set of expressions to choose from, apart from at the final layer where it is a selection of hyperparameters to use in BACON.1. MCTS terminates when the game is over. In chess that would be at checkmate or stalemate, for BACON that would be when each layer of the tree has been searched. Lastly, scoring. In chess, it is +1 for a win, 0 for a draw or -1 for a loss. For BACON it is a function of $\text{MSE}(V_{\text{BACON}}, V + n)$ (referred to as the reward function).

The MCTS algorithm works through 4 steps after being initialised at a given node called a parent node:

1. **SELECTION:** A child node is selected to travel from those available to the parent node (through performing a legal action) based on the score of the child node. The score is weighted by a parameter C . If C is small, the parent node prefers exploitation and picks a child node based on what has returned high scores in the past. If C is larger, the parent prefers exploration and picks a child node which has not been selected much.
2. **EXPANSION:** After the parent chooses a child, this child becomes the parent node. The legal actions are run to determine the new selection of children nodes.
3. **SIMULATION:** This process is repeated until a full game has been simulated. The result of the game is given a score by the reward function.
4. **BACKPROPAGATION:** This score is backpropagated through all the nodes selected to the initial parent node, altering their S_i value.

The purpose of the MCTS is to find the node with the best score after running this process multiple times. It implies this node is the correct one to move to, to maximise winning. In BACON this translates to picking the best expression in a layer that maximises the score. Under an assumption that the best expression in earlier layers will always minimise the MSE (and maximise the score), MCTS is a logical continuation that can search the possible expressions generated by the sets and correctly pick the best expression at each layer. A visualisation of this on the Ideal Gas Law is in Figure 14.

My reward function takes in multiple factors, such as how many variables make up the expression, how quick it took BACON to output the expression as well as the MSE. It is not a continuous function

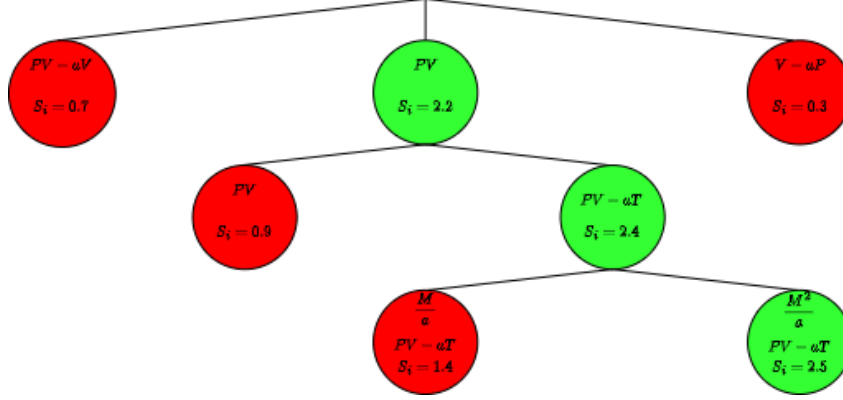


Figure 14: The graphical representation of the MCTS on the low C for the Ideal Gas Law in an environment with 5% noise. The scores at each stage are demonstrated with S_i , with the child node selected in green. Note the green node's S_i increases through the layers as MCTS prunes the worst nodes. For the first two layers the action is picking an equation from those discovered by the sets at that layer. The bottom layer is determined by applying a combination of the following hyperparameters $\epsilon, \delta \in \{(0.01, 0.1), (0.05, 0.5), (0.01, 0), (0, 0.1)\}$ and gaining their respective equations.

and had to be hardcoded for each equation that the MCTS was implemented on. It's hard to determine a general function due to MSE differences varying on different scales between equations.

BACON.M is used to refer to the model that comes from this conjunction with MCTS. It uses BACON.1 at the Space of Laws level, and MCTS to repeatedly traverse the Space of Data.

G.2 Application

Due to time constraints, the system was only adapted for two equations: the Ideal Gas Law and Black's Law. It is simple to adapt for Ohm's Law, however from the previous section it seems likely to perform as the Ideal Gas Law does.

G.2.1 The Ideal Gas Law

The Ideal Gas Law could handle up to 4% noise. MCTS is done on 5% noise at a low C - prioritising exploitation. MCTS outputs BACON.M expression:

$$V_{\text{BACON.M}} = \frac{M(0.502MT + 273.885)}{P} \quad (25)$$

This occurs from the final layer deducing $\frac{M^2}{a}$ as a constant rather than $\frac{M}{a}$. The MSE to $V + n$ is 2.9×10^{-4} and to V is 2.0×10^{-4} . Recalling the graph from Figure 1, this is inline with the predictions for extrapolating to 5% noise in the correct form. It displays that, at this noise level, this form and the correct form share the same MSE to V . *Hence, when there is this much noise, it is not possible to distinguish between the true expression and BACON.M's wrong prediction by comparing to noiseless V .* If this did happen in the real-world, dimensional analysis could be used to show BACON.M's prediction is not valid.

Also, the expression found in MCTS with $\frac{M}{a}$ was

$$V_{\text{BACON.M}} = \frac{M(0.870T + 549.756)}{P} \quad (26)$$

with MSE of 0.026 to $V + n$ and 0.159 to V . The 273 factor is lost - *presenting this level of noise defeats BACON.M*. The coefficients are not related to MCTS, suggesting the process isn't the problem and rather BACON just can't correctly ascertain the best coefficients in 5% noise. This is an upper-bound on what BACON's simplistic mechanisms can handle based on noise not averaging out. PySR on 5% noise finds

$$V_{\text{PySR}} = \frac{M(T + 269.11)}{P} \quad (27)$$

with 1.4×10^{-4} MSE to $V + n$ and 3.2×10^{-5} to V . PySR does not have the drop in performance BACON.M does when noise increases.

G.2.2 Black’s Law

The Ideal Gas Law failed at 5% noise, however the technique may still be valid at less. Reducing the noise for Black’s Law to 2% which failed in subsection F.6, the initial node of $T_f - aT_2$ is selected correctly, but the expression quickly derails into:

$$T_{f\text{BACON.M}} = \frac{(28.544M_2M_1 + 0.940M_2) \left(\frac{M_1}{M_2^3} + \frac{M_1M_2}{M_2^4} \right)^{\frac{1}{3}} - 0.258M_1 - 0.242M_2 + 0.001T_1T_2}{M_1 + 0.940M_2} \quad (28)$$

The MSE loss is 7.704 to $T_f + n$ but 1366 to T_f . The best attempts using BACON.7 is¹²

$$T_{f\text{BACON.7}} = \frac{0.623(0.171M_1^2M_2T_2 - 0.603M_1^2T_1 - 0.553M_1M_2T_1 - M_1M_2T_2 - 0.615M_2^2T_2)}{0.111M_1^3 + 0.102M_1^2M_2 - 0.653M_1^2 - M_1M_2 - 0.368M_2^2} \quad (29)$$

with MSE loss of 13.3 to $T_f + n$ and 12.3 to T_f . *This displays that MCTS overfits on the noisy data.* If instead it was just being defeated by noise we’d expect equation 29 to have higher MSE to T_f than 12.3 which is inline with the exponential increase from Figure 13.

The reason is likely the reward function. Scoring smaller MSEs too high rewards fitting the equation to $T_f + n$ - overfitting. From equation 29, which is the best function we’ve found for fitting T_f , it has an MSE twice equation 28 to $T_f + n$. This is large enough to make me think that other equations can be found that have the same MSE to $T_f + n$ but higher to T_f . The reward function then wouldn’t be able to differentiate based on MSE score, leading to a question of what factors it should consider. Dimensional analysis is a possible direction, but in general *this suggests that MCTS is not the correct adaptation for BACON as the right reward function is very hard to find, if it does exist at all.*

G.3 Discussion

As noise increases, the uncertainty in the true form of the equation becomes a factor. The equations best formed by BACON.M can cut out significant noise, and resemble the true form, whilst having extra factors. This is seen in equation 25. This displays a limitation in the noise BACON can handle. PySR’s attempts at the same problem show that finding the correct form is possible but BACON.M is not powerful enough to get there. Exterior mechanisms to BACON.M would have to be developed such as in a preprocessing stage to remove excess noise.

Whilst the actual implementation between MCTS and BACON was successful, MCTS – and likely any game tree search – is not the solution to improving BACON’s power. It effectively becomes a PySR-like attempt at overfitting on the data, but without the accuracy due to the difficulties in finding a reward function. It is a complex task as what is needed varies between noise and dataset size (seen by different reward functions between the Ideal Gas Law and Black’s Law).

Lastly, MCTS approaches lacks speed. Whilst the Ideal Gas Law took ~ 13 seconds, comparable with PySR, Black’s Law took 700 seconds due to the additional layer and the branching factor associated with the possible equations in a noisy environment. It will not scale well to larger datasets. PySR took ~ 12 seconds to exactly find Black’s Law.

H Related work

H.1 BACON inspired heuristic based search

BACON kick-started a wave of Classical AI in the 1980s exploring CSD. There were multiple projects inspired by Langley using similar heuristic based designs. Two of the most popular were *Fahrenheit* [7] and *Abacus* [5]. *Fahrenheit* makes minor improvements to BACON such as being able to handle irrelevant variables and re-ordering the search algorithm to deal with more terms (covered in limitations in subsection D.3).

¹²Found under args/black/black20.json

Abacus is a more interesting project, outlining similar concerns to mine about BACON. These include dealing with noise, and specifically the problem of approaches such as BACON.6 which require information about the search space (seen in section C). Their solution is threefold. They introduce a more general approach to the search space, such as allowing variables X and Y to have invariant $Y = X^2$ for $X < 3$, but $Y = 12 - X$ for $X \geq 3$. This lets multiple invariants be found within a dataset - something BACON cannot do. They also develop the concept of a proportionality graph search. This uses ideas from graph theory to cycle through the possible variable relationships in multivariable systems. Lastly, they reduce the search space for the invariants by ignoring relationships which can't occur through dimensional analysis.

Fahrenheit's re-ordering of the search algorithm, and Abacus's use of dimensional analysis give fascinating concepts that may be applied to later versions of BACON.7. However, neither deal with noise. One of Fahrenheit's main contributions to deal with irrelevant variables is solved by BACON.7, whilst Abacus yields an approach irreconcilable with BACON. It is also not able to solve a more versatile set of equations than BACON. Hence there is no justification to using either of these projects rather than BACON in attempting to make a noise-resilient Classical AI system. Lastly, no open-source implementations of either means it is hard to meticulously examine for strengths and weaknesses. It is also too time-consuming to make a coded implementation.

The era of Classical AI ends with SDS[12]. It is a mathematically rigorous approach to the same problem, stemming from two postulates by physicist Edgar Buckingham in 1914 that govern the relationship between complete equations and their associated variables scale-type. In other words, given a set of variables and knowledge about how they scale relative to other variables available, a set of possible equations can be formed. This set can be pruned based on the data (using bivariate statistical tests) until only certain equations are possible. The entire algorithm both reduces the complexity found in Langley's BACON, whilst increasing the depth of equations possible. Under the same initial assumptions as BACON, for dependent variable X from a 17 variable circuit meter, the following equations is discoverable:

$$\left(\frac{R_3 h_{fe_2}}{R_3 h_{fe_2} + h_{ie_2}} \frac{R_2 h_{fe_1}}{R_2 h_{fe_1} + h_{ie_1}} \frac{rL^2}{rL^2 + R_1} \right) V - \frac{Q}{C} - \frac{Kh_{ie_3}X}{Bh_{fe_3}} = 0 \quad (30)$$

Whilst this isn't possible in BACON, it displays an increase in flexibility rather than an explicit ability to deal with more noise. On the latter, their best contribution is a 4% relative standard deviation of the dependent variable they claim is solvable through their method - though not backed up in the paper. This becomes the baseline we would like to get BACON to solve. SDS is limited by both the structure the algorithm has to follow and the understandability it lacks - one of the key aspects of using BACON and why it was chosen as the Classical AI approach. Likewise to Fahrenheit and Abacus, there is no open-source project to gain accurate comparisons to our BACON.

None of the Classical AI competitors to BACON share its level of understandability whilst also being more noise-resilient. BACON is the most straightforward algorithm to adapt based on the detail from Langley's textbook.

H.2 Modern neural network approaches

Modern approaches are based around using neural networks. ConservNet proposed by [6] trains a deep feed-forward neural network using their novel noise-variance loss function:

$$\mathcal{L} = \sum_i \text{Var}(F_\theta(\mathbf{x}_{ij})) + |Q - \text{Var}(F_\theta(\mathbf{x}_{ij} + \epsilon_{ij}))|. \quad (31)$$

F_θ represents the conserved function, the left term in the sum represents reducing the intra-group variance of all data that maps to the same value via the function. The right term is the allowed perturbations of the invariant when noise is added to the system - represented by Gaussian noise ϵ_{ij} . Combined it leads to a powerful dynamic where F_θ can't converge trivially, forcing invariants to be found in systems with well-defined Hamiltonians (a measure of total energy in the system). Q controls the scale of the variation. The program is more powerful than BACON. This can be seen in their ability to find the conserved quantity C in environments as varied as real, noisy double pendulum data:

$$C = L_1^2(m_1 + m_2)\omega^2 + m_2 L_2^2 \omega^2 + 2m_1 m_2 L_1 L_2 \omega_1 \omega_2 \cos(\theta_1 - \theta_2) \quad (32)$$

$$- 2gL_1(m_1 + m_2) \cos(\theta_1) - 2gm_2 L_2 \cos(\theta_2) \quad (33)$$

BACON cannot do this.

Similar to SDS, ConservNet exploits a more mathematical approach to improve results. Moreover, ConservNet proves how its neural model converges to appropriate answers and demonstrates substantial experiments dealing with greater noise than any effort we produce with BACON .7. It can also deal with multiple dependent variables which BACON cannot.

Its downsides lie in complexity and explainability. The loss function is novel, but when overlayed on a black-box neural network the meaning is reduced - further shown by the multiple settings needed for them to test to find the optimal hyperparameters for their best model. Additionally, with training times up to several hours, it goes against the principles of the Classical AI programs we are enacting in this project. An interesting implementation would have been putting the trained black-box model at the Space of Laws level; however, due to time constraints this was not possible - though it was done with PySR.

PySR [3] uses symbolic regression to find invariants. Symbolic regression works by imagining the search space as possible mathematical expressions. The expressions are narrowed down to the most accurate and often have a bound on complexity and simplicity in the final result. The latter often manifests itself by reducing constants from floats to their nearest integers. A hypothetical example is when considering Black's law:

$$T_f = \frac{M_1 T_1}{M_1 + M_2} + \frac{M_2 T_2}{M_1 + M_2} \quad (34)$$

With noisy data for independent variables M_1, M_2, T_1 and T_2 , dependent T_f , BACON may find

$$T_{f_{\text{BACON}}} = 1.002 \frac{M_1 T_1}{M_1 + M_2} + 0.844 \frac{M_2 T_2}{M_1 + M_2} \quad (35)$$

whilst PySR finds the exact form – equation 34 – as it assumes the coefficients are 1. PySR is also a much more versatile and powerful program with customisability and the ability to run in large and complex search spaces where the only bound is the computational power. As such, it can handle invariants such as in differential equations and discontinuous environments. BACON cannot compete.

However there are downsides. As it is trained on a deep neural network [4] there is no interface or explainability giving a reason for its outputs. This reduces the trustworthiness of the model.

PySR is used in this project as the main DNN comparison. The reasons why can be seen through two key components in a standard PySR implementation in Figure 15.

```

1 # ml_methods/symbolic_regression.py
2
3 model = PySRRegressor(
4     niterations=15,
5     maxsize=30,
6     binary_operators=["+", "*", "/", "-"],
7     extra_sympy_mappings={"inv": lambda x: 1 / x},
8     loss="loss(prediction, target) = (prediction - target)^2",
9     model_selection="accuracy"
10 )

```

Figure 15: The setup for the PySR implementation used throughout this project. Note *binary_operators* restricting what type of equations can be found, and *niterations* determining how long PySR searches for.

The *niterations* purposes matches that of the complexity counter j in BACON .1 (detailed in subsection B). Experimentally, $j = 4$ and *niterations* = 15 made the complexity of the forms outputted by each program similar. Also, the *binary_operators* is set at $["+", "*", "/", "-"]$. This coincides with the operations that BACON .1 can output. In all, it makes the final equations yielded from our BACON .7 and PySR (in subsection F.6) comparable allowing a balanced discussion on their strengths and weaknesses.

ConservNet's interface didn't trivially allow this level of control. ConservNet also does not have substantial documentation which PySR does, making the PySR implementation straightforward.