

Beyond Loss of Control: Geopolitical Dependency as an Underappreciated Risk of Superintelligence for the Global South

Anonymous Author ^{1,*}

¹ Affiliation withheld for anonymous review

* Correspondence: withheld for anonymous review

Abstract

Discourse on the risks of superintelligence has been dominated by scenarios of loss of control: a future system that escapes human oversight and acts against human interests. We argue that this framing, while important, has obscured a nearer-term and already-materialising risk, which is geopolitical rather than agentic in nature, and which falls disproportionately on the Global South. We identify two compounding mechanisms of dependency that arise as countries build critical infrastructure, in healthcare, finance, and education, on top of a small number of frontier large language models. The first is downstream: access to a frontier system can be withdrawn unilaterally by the jurisdiction in which its developer is incorporated, as illustrated by the 2026 United States export control action suspending foreign access to two frontier models. The second is upstream: frontier models encode the political ideologies and epistemologies of the cultures and corpora that produce them, and a small number of countries currently dominate both training data production and model distribution. We argue that these two mechanisms interact, so that nations dependent on a foreign frontier model for both their information environment and their operational infrastructure are exposed twice over. As frontier AI systems approach and eventually cross thresholds associated with superintelligence, we contend that geopolitical control over access and epistemology, rather than the behaviour of the system itself, constitutes the more immediate and more tractable risk, and one that current superintelligence risk discourse has largely neglected. We close with implications for AI governance frameworks aimed at global equity.

Keywords: superintelligence; AI risk; geopolitics; Global South; AI governance; large language models; data colonialism

1. Introduction

Discussions of the risks posed by superintelligence have, since the earliest formal treatments of the topic, centred on the problem of control: the concern that a sufficiently capable artificial system might pursue objectives misaligned with human interests and that, once such a system exceeds human oversight capacity, this misalignment could not be corrected. This framing has been productive. It has motivated technical research on alignment and interpretability, and it has shaped public and policy debate about the pace and governance of frontier AI development.

We argue that this framing, taken on its own, has left a different and already materialising risk underexamined: the risk that control over access to, and the epistemic content of, frontier AI systems becomes an instrument of geopolitical power, with consequences borne disproportionately by countries that had no voice in the decisions that produced

Received:

Revised:

Accepted:

Published:

Copyright: © 2026 by the authors.

Submitted to *Journal Not Specified* for possible open access publication under the terms and conditions of the

[Creative Commons Attribution \(CC BY\)](#) license.

that dependency. This risk does not require a system to act autonomously against human interests. It arises from entirely human decisions, made by the states and companies that control frontier model development and distribution, and it is already visible in the current generation of large language models, well before those systems reach anything resembling the entity-centric notion of superintelligence that has dominated the field [1].

Large language models are increasingly incorporated into everyday infrastructure: countries and companies are building health, finance, and education applications on top of a small number of frontier systems. We contend that this dependency, combined with two distinct and compounding mechanisms of geopolitical control over frontier models, constitutes an underappreciated risk for the Global South, one that superintelligence risk discourse should treat as a first-order concern rather than a peripheral policy footnote. As frontier systems become more capable, and as the infrastructural dependencies we describe deepen, the geopolitical dimension of this risk does not diminish; it compounds.

2. Two Mechanisms of Geopolitical Dependency

We distinguish two mechanisms by which frontier AI systems become instruments of geopolitical control. The first is downstream and concerns access. The second is upstream and concerns the epistemic content of the systems themselves.

2.1. Downstream: Access as a Revocable Privilege

The downstream mechanism is that access to a frontier model can be withdrawn unilaterally, not by the model itself but by the jurisdiction in which its developer is incorporated. The June 2026 United States export control directive suspending Anthropic’s Fable 5 and Mythos 5 models for foreign nationals [2] demonstrates that frontier model access is contingent on the jurisdiction in which a developer is incorporated, not on the needs of the countries that depend on the model. For a country that has built critical applications, in healthcare diagnostics, financial services, or agricultural advisory systems, on top of a single foreign frontier model, sudden withdrawal of access is not a mere inconvenience but a structural shock to the infrastructure the country has come to rely upon.

Countries in the Global South are, in general, less able to absorb these shocks than more developed economies, since they are less likely to possess a comparable domestically controlled alternative to fall back on. This asymmetry means that the same policy instrument, an export control decision made for reasons entirely internal to the exporting jurisdiction, has starkly unequal downstream effects depending on where a country sits in the global distribution of frontier AI capability.

2.2. Upstream: Epistemic Displacement

The second mechanism is upstream and has received comparatively little attention in superintelligence risk discourse. Frontier models encode the political ideologies and epistemologies of the cultures and corpora that produce them, and the United States and China currently dominate both training-data production and model distribution [3]. Where a national information ecosystem is itself shaped by externally produced or state-coordinated media [3], this dependency is compounded, rather than corrected, by the adoption of foreign frontier models. In Africa, for example, the risk is not merely inaccurate output [4] but the gradual displacement of local epistemic traditions by frameworks embedded, often invisibly, in training data produced elsewhere [5].

2.3. Compounding Effects

These two mechanisms interact, as summarised in Table 1. A nation that depends on a foreign frontier model for both its information environment and its operational infrastructure is exposed twice over: its epistemic autonomy is shaped by what the model

was trained on, and its operational continuity is shaped by who controls the model afterward. Neither risk is mitigated by local hosting alone, since hosting model weights confers little sovereignty if the training corpus and the terms of continued access both originate elsewhere [3,6].

Table 1. Two mechanisms of geopolitical dependency in frontier large language models, and their compounded effect on the Global South.

	Upstream (training data)	Downstream (access control)
Mechanism	Composition of pretraining corpus	Export controls, licensing terms
Controlled by	State media apparatus, corpus curators	Model operator’s home jurisdiction
Primary harm	Epistemic displacement	Operational disruption, vendor lock-in
Most exposed	Language communities under-represented in training data	Nations with critical infrastructure built on a single foreign model

3. Implications for Superintelligence Risk Discourse

We do not claim that current frontier large language models constitute superintelligence in any of the senses debated in the literature, whether the entity-centric threshold of Bostrom [1] or emergent, systemic framings that locate superintelligent behaviour in networks of human and artificial agents rather than in any single system [7]. Our claim is narrower and more urgent: the geopolitical mechanisms described in Section 2 do not depend on any future capability threshold being crossed. They are already operative at current capability levels, and they will not disappear as capabilities increase. If anything, as frontier systems become more capable and more deeply embedded in the critical infrastructure of dependent countries, the leverage that upstream and downstream control confer will grow, not shrink.

This has a direct implication for how the superintelligence risk community allocates its attention. A risk discourse organised primarily around the behaviour of a hypothetical future system risks treating the geopolitical dimension of AI risk as a secondary, merely political concern, external to the technical problem of alignment. We suggest instead that geopolitical dependency should be treated as a first-order risk of the same research and policy programme that studies loss of control, since both concern the distribution of power that frontier AI systems confer, whether that power is exercised by a misaligned system or by the states and companies that control access to an aligned one. The historical record offers a relevant caution here: the telegraph was, in its time, a transformative information technology that nonetheless became an instrument of statecraft and imperial control [8]. Without coordinated governance of training data and access, frontier AI systems risk entrenching a new form of dependency, sometimes termed data colonialism, well before any question of machine autonomy arises.

4. Discussion: Toward Prudent and Reciprocal Governance

We are not arguing that frontier AI systems should be withheld from the Global South. Frontier model access is also a genuine driver of innovation, productivity, and job creation, jobs that economies in the Global South need in order to rise above poverty. The argument is instead for prudence and reciprocity: the benefits of access should not be extended on terms that can be revoked without notice, and the safeguards that justify export controls in principle should not become a mechanism of selective restriction.

We therefore suggest that funders, conferences, and research communities concerned with superintelligence risk and global equity treat frontier model governance as a dual problem, encompassing both training-data provenance and access continuity, rather than addressing either in isolation. Multilateral frameworks for model access, analogous to existing arrangements for satellite spectrum and international telecommunications, would reduce the risk that the upstream and downstream vectors continue to reinforce one another. Absent such frameworks, the countries least able to build frontier alternatives of their own will bear the largest share of both forms of dependency, even as they are told that AI represents an unprecedented opportunity for development.

5. Conclusions

Superintelligence risk discourse has, for good reason, devoted substantial attention to the problem of a future system escaping human control. We argue that this attention should not come at the expense of a nearer-term risk: that control over frontier AI systems, exercised through training-data composition and through unilateral access restriction, is becoming an instrument of geopolitical statecraft, with consequences borne disproportionately by the Global South. These two mechanisms, upstream epistemic displacement and downstream access control, compound one another, and neither is resolved by technical alignment work alone. As frontier systems continue to advance toward whatever threshold the field eventually agrees constitutes “superintelligence”, the geopolitical dependencies we describe will not resolve themselves; they require coordinated governance now, encompassing both training-data provenance and access continuity, if the benefits of frontier AI are to be extended to the Global South on terms that are durable rather than revocable.

Funding: This research received no external funding.

Institutional Review Board Statement: Not applicable.

Informed Consent Statement: Not applicable.

Data Availability Statement: Not applicable. No new data were created or analysed in this study.

Conflicts of Interest: The authors declare no conflicts of interest.

References

1. Bostrom, N. *Superintelligence: Paths, Dangers, Strategies*; Oxford University Press: Oxford, 2014.
2. Anthropic. Update on Fable and Mythos Access. Anthropic, 2026.
3. Waight, H.; Yang, E.; Yuan, Y.; Messing, S.; Roberts, M.E.; Stewart, B.M.; Tucker, J.A. State Media Control Influences Large Language Models. *Nature* **2026**. <https://doi.org/10.1038/s41586-026-10506-7>.
4. Moilola, P. Silicon Valley’s vision for global AI is flawed: each country needs its own blueprint. *Nature* **2026**, *654*, 844. <https://doi.org/10.1038/d41586-026-01951-5>.
5. Birhane, A.; Kalluri, P.; Card, D.; Agnew, W.; Dotan, R.; Bao, M. The Values Encoded in Machine Learning Research. In Proceedings of the Proceedings of the ACM Conference on Fairness, Accountability, and Transparency, 2022.
6. Crawford, K. *Atlas of AI*; Yale University Press, 2021.
7. Author. AI and the Future of the Technosphere: A Path Towards Co-Existence with Superintelligence. In Proceedings of the Proceedings of the 1st AAAI Workshop on Post-Singularity Symbiosis, 2025.
8. Headrick, D.R. *The Invisible Weapon*; Oxford University Press, 1991.

Disclaimer/Publisher’s Note: The statements, opinions and data contained in all publications are solely those of the individual author(s) and contributor(s) and not of MDPI and/or the editor(s). MDPI and/or the editor(s) disclaim responsibility for any injury to people or property resulting from any ideas, methods, instructions or products referred to in the content.